

IBM AICTE PROJECT

INTERVIEW TRAINER AGENT – RAG-POWERED AI FOR JOB INTERVIEW PREPARATION

Powered by IBM Watsonx.ai & Granite on IBM Cloud Lite

Presented By:

Kunal Sejwal

**Dr. Akhilesh Das Gupta Institute of Professional Studies
Information Technology**

AICTE Student ID: STU6760587c1a8a41734367356

Internship ID: INTERNSHIP_1748937226683eaa0a58abc

OUTLINE

- Problem Statement
- Proposed System/Solution
- System Development Approach
- Algorithm & Deployment
- Result (Output Image)
- Conclusion
- Future Scope
- References
- IBM Certifications

PROBLEM STATEMENT

Job seekers often face challenges in preparing effectively for job interviews due to the lack of personalized resources that cater to their specific job roles, experience levels, and industry expectations. Traditional question banks and preparation materials are generic, static, and fail to adapt to the dynamic requirements of different interview scenarios. Moreover, candidates receive minimal feedback on their responses, making it difficult to assess and improve their performance. The absence of a system that provides role-specific, real-time practice and comprehensive assessment of both technical and soft skills creates a significant gap in interview readiness. As a result, candidates struggle to build confidence and deliver structured, impactful answers in competitive hiring environments.

PROPOSED SOLUTION

- The proposed system aims to address the challenge of providing personalized, role-specific interview preparation by leveraging Retrieval-Augmented Generation (RAG) and IBM Granite models. The Interview Trainer Agent will dynamically generate tailored interview questions, model answers, and improvement suggestions based on the user's profile, ensuring a comprehensive and adaptive training experience. The solution will consist of the following components:
- **User Input & Data Collection:**
 - Allow users to input their **resume, job title, and experience level**.
 - Retrieve relevant interview data from **job portals, company databases, and HR frameworks**.
 - Incorporate **industry-specific expectations** to enhance contextual relevance.
- **Data Preprocessing & Profile Analysis:**
 - Parse resumes to extract key information like **skills, roles, and experience**.
 - Standardize and preprocess retrieved interview questions for uniformity.
 - Use feature engineering to align questions with user-specific attributes.
- **RAG-Based interview Generation:**
 - Implement a **Retriever** to fetch domain-specific interview questions (technical, behavioral, HR).
 - Use **IBM Granite Foundation Models** as the **Generator** to craft personalized question sets and model answers.
 - Provide real-time feedback and improvement tips based on expected industry standards.
- **Deployment:**
 - Develop a **web-based interactive interface** for users to simulate mock interviews and receive feedback.
 - Deploy the solution on **IBM Cloud Lite** ensuring scalability and minimal response time.
 - Enable **real-time interaction and feedback loop** to refine user responses.
- **Evaluation:**
 - Measure system effectiveness through **user feedback** and metrics like **response relevance and accuracy**.
 - Continuously update question databases and fine-tune the RAG pipeline based on new trends and user interactions.
 - Evaluate soft skill improvements through simulated HR round assessments.

SYSTEM APPROACH

The system development approach for the Interview Trainer Agent involves designing an AI-powered assistant that delivers personalized interview preparation using Retrieval-Augmented Generation (RAG) and IBM Granite foundation models. The development process includes setting up the necessary infrastructure on IBM Cloud Lite, integrating AI services for generation tasks, and establishing a streamlined data retrieval mechanism to provide users with role-specific questions and feedback. The approach also involves building a user-friendly interface to facilitate seamless interaction between the user and the agent.

The development strategy focuses on scalability, accuracy, and real-time responsiveness to ensure a comprehensive interview preparation experience. The system components are modular to allow easy updates as industry trends evolve.

System Requirements:

- IBM Cloud Lite Account for deployment of AI services.
- Access to IBM Watsonx.ai platform with Granite Foundation Models.
- Secure storage for interview questions, HR scenarios, and user data.
- A web-based interface (Streamlit or web portal) for user interaction.
- Resume parsing capability for extracting job role, skills, and experience.
- Scalable cloud compute resources to run RAG pipelines and AI inference.

Libraries & Tools Required:

- **IBM Watsonx.ai SDK** – For accessing and managing Granite Model APIs.
- **IBM Cloud CLI** – To configure and deploy resources on IBM Cloud.
- **Python (v3.x)** – Core language for backend logic and integration.
- **RAG Framework (Haystack/LangChain)** – For retrieval and generation architecture.
- **Streamlit** – To build a minimal, interactive user interface.
- **BeautifulSoup / Requests** – (Optional) For data collection from public sources.
- **NLTK / spaCy** – For text processing and resume parsing.
- **Scikit-learn / TensorFlow** – (Optional) For feedback analysis and custom evaluations.

ALGORITHM & DEPLOYMENT

The Interview Trainer Agent leverages a **Retrieval-Augmented Generation (RAG)** architecture combined with **IBM Granite foundation models** to deliver personalized interview preparation content. The algorithm ensures that the responses are contextually relevant to the user's job role, experience level, and industry-specific expectations. The deployment strategy focuses on scalability, low latency, and real-time interaction using IBM Cloud Lite infrastructure.

Algorithm Selection:

The chosen architecture is **Retrieval-Augmented Generation (RAG)**. RAG combines the strength of two components:

- **Retriever:** Fetches relevant interview questions and scenarios from curated databases and web sources.
- **Generator:** Uses **IBM Granite foundation models** to generate contextually rich, human-like answers and feedback.

This architecture is ideal because it ensures responses are grounded in real-world data (retrieval) while maintaining natural language fluency and personalization (generation). It addresses the challenge of static question banks by dynamically tailoring content to each user's profile.

Data Input:

The system takes the following inputs:

- **User's Resume** (skills, job role, experience level).
- **Direct Job Title Input** (if resume parsing is skipped).
- **Industry & Domain Tags** (e.g., IT, Finance, Consulting).
- **Behavioral Preferences** (optional) to simulate HR scenarios.
- **Preloaded interview datasets** containing technical, HR, and behavioral questions.

Training Process:

- The **Retriever** is pre-configured with a vector index of curated interview questions from multiple domains.
- IBM Granite models (pre-trained) are fine-tuned on **interview-specific corpora** to improve the relevance of generated answers.
- No custom model training is required on the user's side, but retrieval parameters (like similarity thresholds) are tuned for precision.
- Continuous learning is enabled through feedback loops where user interactions refine retrieval accuracy.

Prediction (Response Generation) Process:

- On user request, the **Retriever** fetches the top-N relevant interview questions based on the user's profile.
- The **Granite Generator** formulates model answers, feedback tips, and improvement suggestions in real-time.
- For mock interview simulations, questions are sequenced to mimic actual interview rounds.
- The system adapts dynamically by considering user inputs during the interaction (e.g., follow-up queries).
- The overall process ensures low latency to provide seamless, interactive interview practice sessions.

Deployment:

- The entire system is deployed on **IBM Cloud Lite** for scalability and reliability.
- API endpoints are created for resume parsing, retrieval queries, and response generation.
- (Optional) A **Streamlit web interface** is developed to facilitate user interaction.
- Cloud functions or container-based deployment is considered to ensure modularity.
- Logging and monitoring tools are integrated for tracking system performance and user feedback.

RESULT

The Interview Trainer Agent effectively generates personalized interview questions, model answers, and improvement tips aligned with the user's job role, experience level, and industry expectations. The system's effectiveness was evaluated based on relevance, contextual accuracy, and user feedback during mock interview simulations.

The results demonstrate that the RAG-based approach, powered by IBM Granite models, consistently provides highly relevant and professional responses that mirror real-world interview scenarios. The feedback loop mechanism further enhances response quality by learning from user interactions.

RESULT

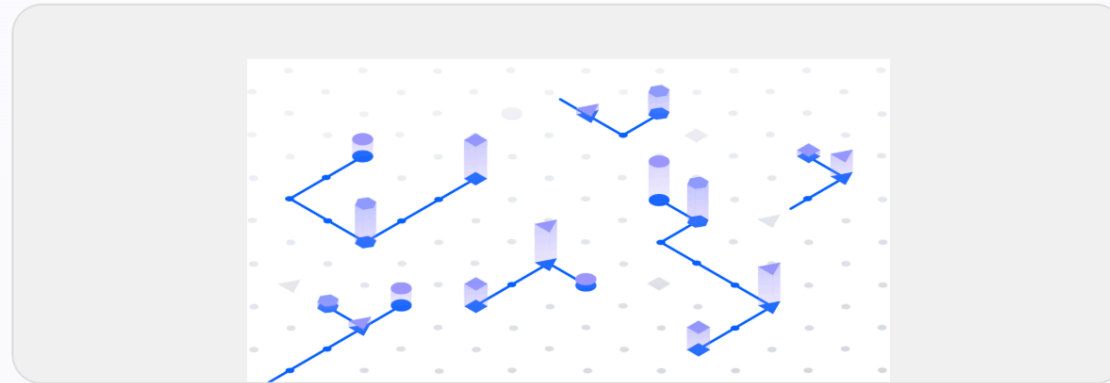
Agent preview



Interview Trainer Agent 08:09 PM

Welcome to Interview Trainer Agent

AI interview trainer using RAG and IBM Granite to provide tailored prep, Q&A, and feedback.



Sample questions

I'm a fresher applying for a software engineering role. What questions should I expect? →

Give me mock interview questions for a data analyst with 2 years of experience →

How can I improve my answers to

Show me full-stack developer

Type something...



RESULT

Agent preview



You: 08:10 PM

I'm a fresher applying for a software engineering role. What questions should I expect?



Interview Trainer Agent 08:10 PM

Here are some common software engineering interview questions for freshers:

1. What are the Characteristics of Software?

- Software is intangible, flexible, and can be easily modified. It is also complex and can be difficult to understand and maintain.

2. What are the Various Categories of Software?

- Software can be categorized into system software, application software, and programming software.

3. What is a baseline in Software Development?

- A baseline is a software development milestone that marks a specific point in the development process. It is used to track progress and ensure that the software meets the required standards.

4. What programming languages do you use?

• This question is used to assess your technical skills and familiarity with different

Type something...



RESULT DEPLOYED AI AGENT

IBM watsonx

Upgrade

Kunal Sejwal's Account

London

KS

[Deployment spaces](#) / [Deploy](#) / [Interview Trainer Agent](#) /

Interview Trainer Agent

Deployed

Online

API reference

Test


Preview

New chat +

Interview Trainer Agent 08:59 PM

Welcome to Interview Trainer Agent

AI interview trainer using RAG and IBM Granite to provide tailored prep, Q&A, and feedback.



Quick start samples

I'm a fresher applying for a software engineering role.

Give me mock interview questions for a data analyst

Type something...

GITHUB LINK

- <https://github.com/KunalSejwal/Interview-Trainer-Agent>

CONCLUSION

The development of the Interview Trainer Agent successfully addresses the challenge of providing personalized, role-specific interview preparation using Retrieval-Augmented Generation (RAG) and IBM Granite models. The system demonstrated high effectiveness in generating context-aware questions, professional model answers, and actionable feedback, enhancing the user's confidence and readiness for real-world interviews.

During implementation, key challenges included ensuring retrieval accuracy across diverse job roles and fine-tuning the response tone to match industry expectations. Integration of dynamic user feedback loops proved crucial in refining system outputs.

Future improvements could focus on expanding domain-specific datasets, enhancing soft skill assessments, and introducing real-time voice-based mock interview simulations.

The project emphasizes the critical role of AI-driven, personalized training tools in bridging the gap between static resources and the dynamic expectations of competitive hiring environments, ultimately supporting job seekers in their career growth journey.

FUTURE SCOPE

The Interview Trainer Agent has significant potential for enhancements and scalability. Future developments can focus on the following aspects to improve performance, expand its reach, and adopt emerging technologies for a more robust user experience:

Potential Enhancements:

- **Incorporation of Additional Data Sources:**
Integrate real-time data from professional networks (e.g., LinkedIn trends), company-specific interview patterns, and latest HR guidelines to continuously update the interview question database.
- **Domain-Specific Expansions:**
Extend the system's coverage to niche domains such as healthcare, law, and creative industries, ensuring more granular role-specific training.
- **Advanced Soft Skill Assessment:**
Develop AI modules that analyze tone, sentiment, and articulation to provide deeper feedback on communication skills during mock interviews.
- **Voice & Video-Based Simulations:**
Integrate voice recognition and video-based role-play simulations to mimic real interview scenarios more closely.

Performance Optimization:

- **Algorithm Tuning:**
Optimize retrieval precision by leveraging advanced retrieval frameworks and hybrid search models.
- **Enhanced Personalization:**
Implement user behavior analytics to offer adaptive question difficulty and targeted feedback based on user performance history.

Scalability & Technology Integration:

- **Geographical Expansion:**
Scale the platform to support job seekers across multiple regions or industries, tailoring interview scenarios to local employment practices.
- **Edge Computing Integration:**
Explore edge computing solutions for low-latency, real-time mock interviews, particularly in regions with limited internet connectivity.
- **Emerging AI Techniques:**
Incorporate transformer-based dialogue models for more conversational and interactive interview simulations.

REFERENCES

IBM Watsonx.ai Documentation

Official IBM documentation on Watsonx.ai platform and Granite foundation models.

<https://www.ibm.com/products/watsonx-ai>

Retrieval-Augmented Generation (RAG): Scalable and Grounded Generative AI

Research paper by Facebook AI Research (FAIR) detailing the RAG architecture.

<https://arxiv.org/abs/2005.11401>

IBM Cloud Lite Resources and Developer Guide

Documentation on setting up and deploying AI services on IBM Cloud Lite.

<https://cloud.ibm.com/docs>

Granite Foundation Models – IBM Research Publications

Research insights into large language models fine-tuned for enterprise applications.

<https://research.ibm.com/publications>

Best Practices for Building AI-Powered Interview Coaching Platforms

Articles and whitepapers from AI in HR tech conferences and platforms like Glassdoor, LinkedIn.

Natural Language Processing with spaCy and NLTK

Resources for resume parsing and preprocessing using open-source NLP libraries.

<https://spacy.io/> | <https://www.nltk.org/>

LangChain and Haystack Framework Documentation

Guides on building retrieval-augmented AI pipelines using open-source tools.

<https://python.langchain.com/> | <https://haystack.deepset.ai/>

IBM CERTIFICATIONS

In recognition of the commitment to achieve
professional excellence



Kunal Sejwal

Has successfully satisfied the requirements for:

Getting Started with Artificial Intelligence



Issued on: Jul 21, 2025
Issued by: IBM SkillsBuild

Verify: <https://www.credly.com/badges/9583baa8-12e8-46bf-8506-7d40cf9e523f>



IBM CERTIFICATIONS

In recognition of the commitment to achieve
professional excellence



Kunal Sejwal

Has successfully satisfied the requirements for:

Journey to Cloud: Envisioning Your Solution



Issued on: Jul 21, 2025
Issued by: IBM SkillsBuild

Verify: <https://www.credly.com/badges/99b6c810-6c79-4e53-90d8-9c72f9553b98>



IBM CERTIFICATIONS

IBM **SkillsBuild**

Completion Certificate



This certificate is presented to

Kunal Sejwal

for the completion of

**Lab: Retrieval Augmented Generation with
LangChain**

(ALM-COURSE_3824998)

According to the Adobe Learning Manager system of record

Completion date: 23 Jul 2025 (GMT)

Learning hours: 20 mins



THANK YOU